

# Reddit SaaS Post Scraper Documentation

Optimized for Supervised Machine Learning - Real Data Only

Version 2.0 - Comment Filtering Approach

## Table of Contents

1. Overview
2. System Architecture
3. Key Features
4. Critical Innovation: Comment Filtering
5. Installation & Setup
6. Configuration Guide
7. Data Collection Process
8. NaN Handling Strategy
9. Weight Validation System
10. Output Files & Structure
11. Troubleshooting
12. Best Practices

## 1. Overview

### Purpose

This scraper collects Reddit posts from entrepreneurship and startup-related subreddits, performs sentiment analysis on posts and comments, and creates a high-quality labeled dataset optimized for supervised machine learning models.

### Critical Innovation (v2.0)

**Comment Filtering Strategy** - Only collects posts with minimum 3 comments to ensure ALL posts have valid comment sentiment scores. This eliminates the NaN problem and provides natural variance in comment sentiment without artificial noise.

## Key Guarantee

**No Zero Weights** - All validation parameters (post sentiment, comment sentiment, upvote ratio, post recency) are guaranteed to contribute to the final validation score with real data variance.

## Target Use Case

Training supervised ML models to classify SaaS ideas as "good", "neutral", or "bad" based on:

- Post content sentiment
- Public feedback (comment sentiment) - **Always present**
- Community engagement (upvote ratio)
- Temporal relevance (post recency)

## 2. System Architecture

### Components

#### Core Modules

1. **reddit\_scraper.py** - Main scraping engine with comment filtering
2. **helpers.py** - Utility functions for data processing
3. **weight\_validator.py** - Automated weight optimization and labeling
4. **config.yaml** - Configuration file for all parameters

#### External Dependencies

- **PRAW** - Reddit API wrapper
- **Transformers** - Hugging Face sentiment analysis model
- **Pandas** - Data manipulation
- **NumPy** - Numerical operations

#### Data Flow

```
Reddit API → Post Filtering (min comments) → Sentiment Analysis →  
NaN Handling (minimal) → Weight Optimization → Labeling → CSV Export
```

### 3. Key Features

#### 3.1 Multi-Source Data Collection

Rotates through 7 different Reddit post sources:

1. New Posts
2. Hot/Trending Posts
3. Top Posts (Week)
4. Top Posts (Month)
5. Top Posts (Year)
6. Controversial Posts (Week)
7. Controversial Posts (Month)

**Benefit:** Maximum diversity in dataset, prevents sampling bias

#### 3.2 Dual Sentiment Analysis

- **Post Sentiment:** Analyzes title + body text
- **Comment Sentiment:** Analyzes up to 15 newest comments (public feedback)

**Model:** DistilBERT (fine-tuned for sentiment analysis)

#### 3.3 Smart Comment Filtering

- **Skips posts with < 3 comments** (configurable)
- Ensures all collected posts have valid comment sentiment
- Natural variance in comment sentiment (no artificial noise needed)

**This is the key to solving the zero-weight problem**

#### 3.4 Smart Deduplication

- Global tracking of collected post IDs across all batches
- Prevents duplicate posts in dataset
- Automatic retry with fallback sources

#### 3.5 Minimal NaN Handling

- Since posts are filtered, very few NaN values
- Simple median replacement (no noise needed)
- Enforces minimum value of 0.01

### 3.6 Minimum Weight Enforcement

- All parameters guaranteed to have weight  $\geq 0.01$
- Prevents unused features in validation
- Better for supervised ML training

### 3.7 Automated Per-Batch Weight Optimization

- Tests 286 weight combinations per batch
- Finds optimal weights based on label distribution entropy
- Adapts to data characteristics of each batch

## 4. Critical Innovation: Comment Filtering

### The Problem (v1.0)

Collected 100 posts:

- 55 posts with comments → real comment sentiment
- 45 posts with NO comments → NaN comment sentiment

After NaN replacement with mean/median:

- 45 identical values → LOW variance
- Optimizer says "comment sentiment is useless"
- Weight forced to minimum 0.01

Result: Comment sentiment not actually used!

### The Solution (v2.0)

Filter during collection:

- Check if post has  $\geq 3$  comments
- If yes → collect it
- If no → skip it and fetch next post

After collection:

- ALL 100 posts have real comment sentiments
- Natural variance (different posts, different comments)
- NO NaN problem!

Result: Comment sentiment weight = 0.15-0.35 naturally!

## Implementation

```
# In scraper
MIN_COMMENTS_REQUIRED = 3 # From config.yaml

# Skip posts with too few comments
if num_comments < MIN_COMMENTS_REQUIRED:
    total_no_comments += 1
    continue # Don't collect this post

# If comment sentiment still can't be calculated, skip
if len(comment_sentiments) == 0:
    total_no_comments += 1
    continue
```

## Trade-offs

### Pros:

- Real data only (no artificial noise)
- Natural variance in comment sentiment
- Comment sentiment actually used in validation
- Cleaner dataset

### Cons:

- Takes longer to collect 100 posts (may need to check 120-150)
- Misses posts with no comments (but those aren't useful for comment sentiment anyway)

**Verdict:** The pros far outweigh the cons. This is the correct approach.

## 5. Installation & Setup

### Prerequisites

```
Python 3.8+
pip (Python package manager)
Reddit API credentials
```

### Step 1: Install Dependencies

```
pip install praw pandas numpy transformers python-dotenv pyyaml torch
```

## Step 2: Reddit API Setup

1. Go to <https://www.reddit.com/prefs/apps>
2. Create a new application (script type)
3. Note down:
  - Client ID
  - Client Secret
  - User Agent

## Step 3: Environment Variables

Create `.env` file in project root:

```
REDDIT_CLIENT_ID=your_client_id_here
REDDIT_SECRET=your_client_secret_here
REDDIT_USER_AGENT=your_user_agent_here
```

## Step 4: Directory Structure

```
project_root/
├── src/
│   ├── data_collection/
│   │   └── reddit_scraper.py
│   ├── utils/
│   │   ├── helpers.py
│   │   └── weight_validator.py
│   ├── config.yaml
│   ├── .env
│   └── data/
│       └── raw/
│           ├── raw_batch/           # Output CSV files
│           └── raw_batch_report/    # Optimization reports
```

## 6. Configuration Guide

### config.yaml Structure

### Scraper Settings

```
scraper:
  subreddits:
    - Entrepreneur
    - startups
    - indiehackers
  batch_size: 100          # Posts per batch
  max_batches: 10         # Total batches
```

```
max_comments_per_post: 15 # For sentiment analysis

# CRITICAL: Minimum comments required
min_comments_required: 3 # Skip posts with < 3 comments

delay_min: 1.0          # Min delay between requests
delay_max: 3.0          # Max delay between requests
```

**Key Parameter:** min\_comments\_required

- Default: 3
- Recommended: 3-5
- Higher = more selective, longer collection time
- Lower = faster collection, potentially more NaN

## NaN Handling

```
nan_handling:
  method: 'median' # Simple median (no noise needed!)
  min_value: 0.01  # Minimum value enforcement
```

**Note:** With comment filtering, NaN is rare, so simple median works fine.

## Validation Thresholds

```
validation_thresholds:
  good: 70      # Score >= 70 → "good"
  neutral: 40   # 40 <= Score < 70 → "neutral"
              # Score < 40 → "bad"
```

## Weight Validator

```
weight_validator:
  step_size: 0.1    # Weight granularity
  min_weight: 0.01  # Minimum weight per parameter
```

## 7. Data Collection Process

### Phase 1: Post Extraction with Filtering

For each batch:

1. **Select Source:** Rotates through 7 sources based on batch number
2. **Fetch Posts:** Starts fetching from selected source
3. **Age Filter:** Checks if post age  $\leq$  180 days

4. **Deduplication Check:** Checks if post ID already collected
5. **Comment Filter: NEW** - Checks if post has  $\geq$  min\_comments\_required
6. **Extract Metadata:** If all filters pass

**Key Difference from v1.0:** Step 5 is new and critical

## Phase 2: Sentiment Analysis

### Post Sentiment

```
full_text = title + " " + body
post_sentiment = sentiment_model(full_text)
# Returns: 0.0 (negative) to 1.0 (positive)
# Returns: NaN if text too short
```

### Comment Sentiment

```
# Fetch newest 15 comments (already filtered for min count)
comments = sorted_by_newest[:15]

# Analyze each comment
comment_sentiments = [sentiment_model(c.body) for c in comments]

# If no valid sentiments after analysis, SKIP THIS POST
if len(comment_sentiments) == 0:
    skip_post()
    continue

# Calculate average (guaranteed to have values)
avg_comment_sentiment = mean(comment_sentiments)
# Returns: Real value, very rarely NaN
```

## Phase 3: Feature Calculation

```
features = {
    'post_sentiment': 0.0 - 1.0,
    'avg_comment_sentiment': 0.0 - 1.0, # Always valid!
    'upvote_ratio': 0.0 - 1.0,
    'post_recency': exp(-days_old / 30)
}
```

## Collection Statistics

### Expected:

- Posts examined: ~120-150 per batch
- Posts collected: 100 per batch



- Posts skipped (no comments): ~20-50 per batch
- Collection time: 8-12 minutes per batch (slightly longer)

## 8. NaN Handling Strategy

### Why NaN Values Are Minimal

With comment filtering:

- **Post Sentiment NaN:** Rare (only if text too short)
- **Comment Sentiment NaN:** Very rare (posts filtered, comments analyzed)
- **Upvote Ratio NaN:** Rare (Reddit API issue)

### Handling Method: Simple Median

```
# For the few NaN that remain
median_value = df['feature'].median()
df['feature'].fillna(median_value)
```

### Why median instead of mean+noise:

- NaN count is very low (< 5%)
- Natural variance already high
- No need for artificial noise
- Simpler and more honest

### Comparison

| Method               | NaN Count  | Approach             | Comment Weight   |
|----------------------|------------|----------------------|------------------|
| v1.0 (no filter)     | ~45%       | mean+noise           | 0.01 (failed)    |
| <b>v2.0 (filter)</b> | <b>~2%</b> | <b>simple median</b> | <b>0.15-0.35</b> |

## 9. Weight Validation System

### Weight Components

```
weights = [w1, w2, w3, w4]  # Sum = 1.0

where:
  w1 = post_sentiment weight
  w2 = avg_comment_sentiment weight
```

```
w3 = upvote_ratio weight
w4 = post_recency weight
```

## Validation Score Calculation

```
score = (
    w1 * post_sentiment +
    w2 * avg_comment_sentiment +
    w3 * upvote_ratio +
    w4 * post_recency
) * 100

# Range: 0 - 100
```

## Label Assignment

```
if score >= 70:
    label = "good"
elif score >= 40:
    label = "neutral"
else:
    label = "bad"
```

## Expected Weight Distribution (v2.0)

With real comment sentiment variance:

```
Typical weights:
- Post Sentiment:      0.25 - 0.40
- Comment Sentiment:   0.15 - 0.35 ← Actually used!
- Upvote Ratio:        0.15 - 0.30
- Post Recency:        0.15 - 0.30
```

## Example Output

```
Best Weights (All Parameters Used):
- Post Sentiment:      0.32 [USED]
- Comment Sentiment (avg): 0.24 [USED] ← Real weight!
- Upvote Ratio:        0.22 [USED]
- Post Recency:        0.22 [USED]
- Total: 1.00

Comment Sentiment Variance: 0.084523 ← High variance!
```

## 10. Output Files & Structure

### 10.1 CSV Files

**Location:** data/raw/raw\_batch/

**Filename:** reddit\_data\_batch\_{N}.csv

**Columns:**

|                       |   |
|-----------------------|---|
| post_id               | - Unique Reddit post ID                       |
| subreddit             | - Source subreddit                            |
| title                 | - Post title                                  |
| text                  | - Title + body combined                       |
| author                | - Reddit username                             |
| created_utc           | - Timestamp (ISO format)                      |
| num_comments          | - Number of comments (always >= min_required) |
| upvotes               | - Score from Reddit                           |
| upvote_ratio          | - Ratio of upvotes (0-1)                      |
| post_age_days         | - Age in days                                 |
| post_sentiment        | - Sentiment score (0-1)                       |
| avg_comment_sentiment | - Avg comment sentiment (0-1, rarely NaN)     |
| post_recency          | - Recency weight (0-1)                        |
| validation_score      | - Final score (0-100)                         |
| label                 | - good/neutral/bad                            |
| source_url            | - Reddit permalink                            |
| source_type           | - Data source                                 |

### 10.2 Report Files

**Location:** data/raw/raw\_batch\_report/

**Filename:** reddit\_data\_batch\_{N}\_report.txt

**Contents** (v2.0 format):

```
BATCH 1 - WEIGHT OPTIMIZATION RESULTS
=====

Source: NEW_POSTS
Posts skipped (no comments): 38

Optimization Accuracy: 98.45

Best Weights:
- Post Sentiment:          0.32
- Comment Sentiment (avg): 0.24 ← Real weight!
- Upvote Ratio:           0.22
- Post Recency:           0.22

Weights String: 0.32,0.24,0.22,0.22

Comment Sentiment Variance: 0.084523 ← High!
```

=====

## 11. Troubleshooting

### Issue 1: Collection Takes Too Long

**Symptom:** Taking > 15 minutes per batch

**Cause:** High `min_comments_required` or subreddits with few commented posts

**Solution:**

- Lower `min_comments_required` to 2-3
- Add more active subreddits to the list
- Increase `batch_size` slightly to compensate

### Issue 2: Comment Sentiment Weight Still Low

**Symptom:** Weight = 0.01-0.05 even with filtering

**Cause:** Comments are too similar in sentiment (all positive or all negative)

**Solution:**

- Add controversial subreddits (more varied opinions)
- Increase `max_comments_per_post` to 20
- Check if sentiment model is working correctly

### Issue 3: Too Many Posts Skipped

**Symptom:** "Posts skipped (no comments): 80+"

**Cause:** Source has many low-engagement posts

**Solution:**

- Use "Hot" or "Top" sources instead of "New"
- Lower `min_comments_required` to 2
- Focus on more active subreddits

### Issue 4: Comment Sentiment Variance Check Shows Low Value

**Symptom:** Variance < 0.02

**Solution:**

```
# In scraper output, look for:  
[VARIANCE CHECK]  
    avg_comment_sentiment: 0.015000 ← Too low!  
  
# Then:  
1. Increase min_comments_required to 5-7  
2. Use more diverse subreddits  
3. Check sentiment model is analyzing correctly
```

## Issue 5: NaN in Comment Sentiment

**Symptom:** Despite filtering, some posts still have NaN

**Cause:** Comments exist but all fail sentiment analysis (too short, special chars, etc.)

**Solution:** Already handled - scraper skips these posts

```
if len(comment_sentiments) == 0:  
    total_no_comments += 1  
    continue
```

## 12. Best Practices

### Data Quality

#### 1. Comment Filtering Configuration

```
# Recommended settings  
min_comments_required: 3-5  
  
# Too low (1-2): Some posts may still lack diverse comments  
# Recommended (3-5): Good balance  
# Too high (> 7): Very slow collection, may skip quality posts
```

#### 2. Batch Size

- **Recommended:** 100 posts per batch
- Actual posts examined: ~120-150 with filtering
- Collection time: ~10 minutes per batch

### 3. Number of Batches

- **Recommended:** 10 batches minimum
- With filtering: 1000 high-quality posts
- Total time: ~2 hours for full dataset

### 4. Subreddit Selection

Choose active subreddits with engaged communities:

- Entrepreneur (high volume, many comments)
- startups (quality discussions)
- indiehackers (active builder community)
- microsaas (niche but engaged)

### Comment Filtering

#### Optimal Settings

**For Maximum Quality** (slower):

```
min_comments_required: 5
max_comments_per_post: 20
```

**For Balanced Approach** (recommended):

```
min_comments_required: 3
max_comments_per_post: 15
```

**For Faster Collection:**

```
min_comments_required: 2
max_comments_per_post: 10
```

### Monitoring Collection

Watch for these metrics:

```
> Collecting: 67/100 | No comments skipped: 42
```

If "No comments skipped" is very high:

- Lower min\_comments\_required
- Change to more active source (Hot instead of New)

## Variance Verification

After each batch, check:

```
[VARIANCE CHECK]
post_sentiment: 0.142000 ← Good
avg_comment_sentiment: 0.078000 ← Good (should be > 0.05)
upvote_ratio: 0.039000 ← OK
```

If comment sentiment variance < 0.05:

- Increase min\_comments\_required
- Add more diverse subreddits
- Check sentiment model

## Supervised ML Training

### Feature Usage

With v2.0, all features have real contribution:

- Post Sentiment: Always used (0.25-0.40)
- **Comment Sentiment: Now properly used (0.15-0.35)**
- Upvote Ratio: Always used (0.15-0.30)
- Post Recency: Always used (0.15-0.30)

### Model Training

```
# All 4 features are meaningful
X = df[['post_sentiment', 'avg_comment_sentiment', 'upvote_ratio', 'post_recency']]
y = df['label']

# No need to drop any features!
model.fit(X, y)
```

## Appendix A: v1.0 vs v2.0 Comparison

### Key Differences

| Aspect       | v1.0         | v2.0                    |
|--------------|--------------|-------------------------|
| Collection   | All posts    | Filtered (min comments) |
| Comment NaN  | ~45%         | ~2%                     |
| NaN Handling | mean + noise | simple median           |

| Aspect           | v1.0           | v2.0             |
|------------------|----------------|------------------|
| Comment Variance | Low (0.02)     | High (0.08)      |
| Comment Weight   | 0.01 (min)     | 0.15-0.35 (real) |
| Collection Time  | 5-7 min        | 8-12 min         |
| Data Quality     | Mixed          | High             |
| Approach         | Artificial fix | Natural solution |

## Why v2.0 is Better

1. **Real Data:** No artificial noise, just actual comment sentiments
2. **Natural Variance:** Different posts genuinely have different comment tones
3. **All Features Used:** Comment sentiment actually contributes
4. **Cleaner Dataset:** All posts have the features we care about
5. **ML-Ready:** Better for training supervised models

## When to Use Each

**v1.0** (with noise):

- When you need to collect ALL posts regardless of comments
- When dataset size is more important than quality
- Research/analysis where comment sentiment is not critical

**v2.0** (with filtering):

- **For supervised ML training** (recommended)
- When comment sentiment is important
- When you want clean, high-quality data
- Production use cases

## Appendix B: Complete Workflow

### Step-by-Step Process

#### 1. Setup (One-time)

```
pip install requirements
Configure Reddit API in .env
Set min_comments_required in config.yaml
```

#### 2. Run Scraper



```
python src/data_collection/reddit_scraper.py
```

### 3. Monitor Progress

Watch for:

- Posts skipped (no comments)
- Variance check values
- Weight distribution

### 4. Verify Quality

```
import pandas as pd
df = pd.read_csv('data/raw/raw_batch/reddit_data_batch_1.csv')

# Check NaN count
print(df['avg_comment_sentiment'].isna().sum()) # Should be < 5

# Check variance
print(df['avg_comment_sentiment'].var()) # Should be > 0.05
```

### 5. Combine Batches

```
import glob
csv_files = glob.glob('data/raw/raw_batch/*.csv')
df_combined = pd.concat([pd.read_csv(f) for f in csv_files])
df_combined.to_csv('combined_dataset.csv', index=False)
```

### 6. Train ML Model

```
from sklearn.ensemble import RandomForestClassifier

X = df_combined[['post_sentiment', 'avg_comment_sentiment',
                 'upvote_ratio', 'post_recency']]
y = df_combined['label']

model = RandomForestClassifier()
model.fit(X, y)
```

## Appendix C: Configuration Examples

### Example 1: High Quality, Slower

```
scraper:
  batch_size: 100
  max_batches: 10
  min_comments_required: 5
  max_comments_per_post: 20

# Expected: 1000 posts in ~2.5 hours
```

```
# Quality: Excellent
# Comment weight: 0.25-0.40
```

## Example 2: Balanced (Recommended)

```
scraper:
  batch_size: 100
  max_batches: 10
  min_comments_required: 3
  max_comments_per_post: 15

# Expected: 1000 posts in ~2 hours
# Quality: Very good
# Comment weight: 0.15-0.35
```

## Example 3: Quick Test

```
scraper:
  batch_size: 50
  max_batches: 2
  min_comments_required: 2
  max_comments_per_post: 10

# Expected: 100 posts in ~15 minutes
# Quality: Good
# For testing only
```

## Conclusion

Version 2.0 solves the critical comment sentiment weight problem through a simple but effective approach: **only collect posts that have comments**. This ensures natural variance in comment sentiment without resorting to artificial noise, resulting in a cleaner, higher-quality dataset perfect for supervised ML training.

### Key Improvements in v2.0:

- Comment filtering ensures all posts have valid comment sentiment
- Natural variance (no artificial noise needed)
- Comment sentiment weight 15-35% (vs 1% in v1.0)
- All 4 validation parameters properly used
- Better dataset quality for ML

**Recommendation:** Use v2.0 for all supervised ML applications.

**Document Version:** 2.0

**Last Updated:** October 29, 2025

**Critical Change:** Comment filtering strategy implemented